

Agentic Honeypot API: Complete Build Plan & PRD

Product: SentinelHP - Agentic Honeypot for Scam Detection & Intelligence Extraction

Hackathon: GUVI/HCL AI for Impact - AI for Fraud Detection & User Safety

Timeline: 2 Days with Integrated Testing

Prepared: February 2026

Executive Summary

This document provides a comprehensive plan to build and deploy an Agentic Honeypot API system within 2 days. The system detects scam messages, autonomously engages scammers using AI agents, extracts actionable intelligence (UPI IDs, bank accounts, phishing links), and reports structured results to an evaluation endpoint.

The plan integrates development, testing, and API endpoint validation as required by organizers.

Part 1: Simplified 2-Day Build Plan

Day 1: Core Infrastructure (8-10 hours)

Phase 1A: API Foundation (2-3 hours) - UPDATED FOR GUVI FORMAT

Objective: Working API endpoint with GUVI-compliant schema

Tasks:

1. Set up FastAPI project structure
2. Implement API key authentication middleware (header: x-api-key)
3. Create POST endpoint (e.g., /api/honeypot or /api/message)
4. Add request schema validation for GUVI format
5. Implement basic logging

Request Schema (Pydantic Models):

```
class Message(BaseModel):
    sender: str # "scammer" or "user"
    text: str
    timestamp: str

class ConversationMessage(BaseModel):
    sender: str
    text: str
    timestamp: str

class Metadata(BaseModel):
    channel: str
    language: str
    locale: str

class HoneypotRequest(BaseModel):
    sessionId: str
    message: Message
    conversationHistory: List[ConversationMessage] = []
    metadata: Optional[Metadata] = None

class HoneypotResponse(BaseModel):
    status: str = "success"
    reply: str
```

Deliverables:

- Working API accepting GUVI-format requests
- API key validation functional (`x-api-key` header)
- Response format matching GUVI specs

Testing Checkpoint:

```
curl -X POST https://your-endpoint/api/honeypot
-H "x-api-key: your-secret-key"
-H "Content-Type: application/json"
-d '{
  "sessionId": "test-session-001",
  "message": {
    "sender": "scammer",
    "text": "Your bank account will be blocked today. Verify immediately.",
    "timestamp": "2026-02-02T10:15:30Z"
  },
  "conversationHistory": [],
  "metadata": {
    "channel": "SMS",
    "language": "English",
    "locale": "IN"
  }
}'
```

Expected response:

```
{
  "status": "success",
```

```
"reply": "Why will my account be blocked?"  
}
```

Phase 1B: Database Setup (1-2 hours)

Objective: Session persistence with PostgreSQL

Schema Design:

```
-- Sessions table
```

```
CREATE TABLE sessions (  
    session_id VARCHAR(255) PRIMARY KEY,  
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,  
    updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,  
    status VARCHAR(50),  
    risk_score FLOAT,  
    agent_engaged BOOLEAN DEFAULT FALSE  
)
```

```
-- Messages table
```

```
CREATE TABLE messages (  
    id SERIAL PRIMARY KEY,  
    session_id VARCHAR(255) REFERENCES sessions(session_id),  
    turn_number INTEGER,  
    sender VARCHAR(50),  
    message TEXT,  
    timestamp TIMESTAMP,  
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP  
)
```

```
-- Extracted entities table
```

```
CREATE TABLE extracted_entities (  
    id SERIAL PRIMARY KEY,  
    session_id VARCHAR(255) REFERENCES sessions(session_id),  
    entity_type VARCHAR(50),  
    entity_value TEXT,  
    confidence FLOAT,  
    source_turn INTEGER,  
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP  
)
```

Tasks:

1. Set up PostgreSQL (or SQLite for quick demo)
2. Create database schema
3. Implement SQLAlchemy models
4. Add session management functions

Testing Checkpoint:

- Verify session creation in database

- Test message logging
 - Confirm entity insertion works
-

Phase 1C: Scam Detection Layer (3-4 hours)

Objective: Layered detection system (rules → LLM)

Rule-Based Heuristics:

```
SCAM_KEYWORDS = {  
    'urgency': ['urgent', 'immediately', 'now', 'today', 'blocked', 'suspended'],  
    'authority': ['bank', 'police', 'government', 'RBI', 'income tax', 'courier'],  
    'financial': ['refund', 'cashback', 'prize', 'lottery', 'verify account', 'KYC'],  
    'credentials': ['OTP', 'PIN', 'password', 'CVV', 'card number', 'Aadhaar']  
}
```

```
ENTITY_PATTERNS = {  
    'upi': r'\b[\w.-]+@[^\w]+\b',  
    'phone': r'+91[\d]{10}|\b[6-9]\d{9}\b',  
    'url': r'https://[^\\s]+',  
    'bank_account': r'\b\d{9,18}\b'  
}
```

Detection Pipeline:

1. **Rule Layer:** Score based on keyword presence, entity detection, urgency patterns
2. **Behavioral Layer:** Unknown sender + financial request + urgency = high risk
3. **LLM Fallback:** Google Gemini classification for ambiguous cases

Output Structure:

```
{  
    "risk_score": 0.87,  
    "confidence": "high",  
    "signals": ["urgency_detected", "upi_present", "authority_impersonation"],  
    "should_engage": true  
}
```

Testing Checkpoint:

```
test_messages = [  
    "Your UPI will be blocked. Verify now: fraud@paytm",  
    "Hello, how are you?",  
    "URGENT: KYC update needed. Call +919876543210"  
]
```

Verify correct risk scores and signal detection

Phase 1D: Google ADK Agent Setup (2 hours)

Objective: Working AI agent with persona

System Prompt Template:

You are a human user engaging with someone via SMS/WhatsApp.

PERSONA: {persona_description}

STRICT RULES:

1. Never admit you are an AI
2. Never provide real credentials, bank details, or OTPs
3. Ask clarifying questions when confused
4. Show mild hesitation and request repetition
5. Mirror the language style of the sender
6. Never confront or accuse anyone of scamming

CONVERSATION GOAL:

Understand what the sender wants while appearing cooperative but slightly confused.

MEMORY:

{conversation_history}

Respond naturally in 1-2 sentences.

Persona Examples:

- **Digitally Naive:** "I'm not very good with technology. Which app should I open exactly?"
- **Average User:** "I tried but it shows pending. What should I do next?"
- **Skeptical:** "This seems unusual. Can you confirm you're from the bank?"

Testing Checkpoint:

- Send test message to agent
- Verify human-like response
- Confirm no AI disclosure

Day 2: Intelligence & Integration (8-10 hours)

Phase 2A: Intelligence Extraction Pipeline (2-3 hours)

Objective: High-confidence entity extraction

Extraction Logic:

```
class IntelligenceExtractor:  
    def extract_entities(self, message: str, conversation_history: list) -> dict:  
        entities = {  
            'upi_ids': [],  
            'bank_accounts': [],  
            'phone_numbers': [],  
            'urls': []  
        }  
  
        # Regex extraction  
        entities['upi_ids'] = self.validate_upi(re.findall(UPI_PATTERN,  
message))  
        entities['phone_numbers'] =  
self.validate_phone(re.findall(PHONE_PATTERN, message))  
        entities['urls'] = self.validate_url(re.findall(URL_PATTERN,  
message))  
  
        # Contextual confirmation using LLM  
        confirmed = self.llm_confirm_entities(entities,  
conversation_history)  
  
        # Only return high-confidence (>0.85)  
        return self.filter_by_confidence(confirmed, threshold=0.85)
```

Validation Rules:

| Entity Type | Validation |
|--------------|--|
| UPI ID | Format: user@provider, provider in known list |
| Phone | +91 prefix or 10-digit starting with 6-9 |
| Bank Account | 9-18 digits, checksum if IFSC present |
| URL | Valid domain, not whitelisted legitimate sites |

Table 1: Entity Validation Criteria

Testing Checkpoint:

```
test_message = "Send payment to fraud123@paytm. Call me at +919876543210. Visit  
http://fake-bank.site"  
extracted = extractor.extract_entities(test_message, [])  
assert len(extracted['upi_ids']) == 1  
assert len(extracted['phone_numbers']) == 1  
assert extracted['upi_ids'][0]['confidence'] > 0.85
```

Phase 2B: Multi-Turn Conversation Loop (2 hours)

Objective: Sustained agent engagement with memory

Conversation Controller:

```
class ConversationController:
    MAX_TURNS = 10
    TIMEOUT_SECONDS = 300

    async def handle_turn(self, session_id: str, message: str):
        session = self.db.get_session(session_id)

        # Check termination conditions
        if self.should_terminate(session):
            await self.finalize_session(session_id)
            return

        # Load conversation history
        history = self.db.get_messages(session_id)

        # Generate agent response
        response = await self.agent.generate_response(
            message=message,
            history=history,
            persona=session.persona
        )

        # Extract intelligence from current turn
        entities = self.extractor.extract_entities(message, history)

        # Store turn
        self.db.save_message(session_id, message, response)
        self.db.save_entities(session_id, entities)

        return response

    def should_terminate(self, session) -> bool:
        return (
            session.turn_count >= self.MAX_TURNS or
            session.duration > self.TIMEOUT_SECONDS or
            session.intelligence_confidence >= 0.9 or
            session.safety_override
        )
```

Testing Checkpoint:

- Simulate 5-turn conversation
 - Verify memory persistence
 - Confirm termination conditions work
-

Phase 2C: Evaluation Callback Implementation (1-2 hours) - GUVI COMPLIANT

Objective: Reliable callback to GUVI evaluation endpoint

Callback Endpoint:

<https://hackathon.guvi.in/api/updateHoneyPotFinalResult>

EXACT Callback Payload Structure (CRITICAL - Must match GUVI spec):

```
{  
    "sessionId": "abc123-session-id",  
    "scamDetected": true,  
    "totalMessagesExchanged": 18,  
    "extractedIntelligence": {  
        "bankAccounts": ["XXXX-XXXX-XXXX", "9876543210"],  
        "upiIds": ["scammer@upi", "fraud@paytm"],  
        "phishingLinks": ["http://malicious-link.example", "http://fake-bank.site"],  
        "phoneNumbers": ["+91XXXXXXXXXX", "+919876543210"],  
        "suspiciousKeywords": ["urgent", "verify now", "account blocked", "KYC update"]  
    },  
    "agentNotes": "Scammer used urgency tactics and payment redirection. Impersonated bank authority."  
}
```

Implementation Example:

```
class EvaluationCallback:  
    GUVI_ENDPOINT = "https://hackathon.guvi.in/api/updateHoneyPotFinalResult"  
    MAX_RETRIES = 3  
    RETRY_DELAY = [1, 3, 5]  
  
    @async def send_callback(self, session_id: str):  
        # Build GUVI-compliant payload  
        payload = self.build_guvi_payload(session_id)  
  
        for attempt in range(self.MAX_RETRIES):  
            try:  
                response = await self.http_client.post(  
                    url=self.GUVI_ENDPOINT,  
                    json=payload,  
                    headers={"Content-Type": "application/json"},  
                    timeout=10  
                )  
  
                if response.status_code == 200:  
                    self.db.mark_callback_sent(session_id)  
                    logger.info(f"GUVI callback successful: {session_id}")  
                    return  
  
            except Exception as e:  
                logger.error(f"GUVI callback attempt {attempt+1} failed:  
{e}")
```

```

{e}")
        await asyncio.sleep(self.RETRY_DELAY[attempt])

    self.db.mark_callback_failed(session_id)

def build_guvi_payload(self, session_id: str) -> dict:
    session = self.db.get_session(session_id)
    entities = self.db.get_entities(session_id)
    messages = self.db.get_messages(session_id)

    return {
        "sessionId": session_id,
        "scamDetected": session.scam_detected,
        "totalMessagesExchanged": len(messages),
        "extractedIntelligence": {
            "bankAccounts": [e['value'] for e in entities if e['type'] ==
== 'bank_account'],
            "upiIds": [e['value'] for e in entities if e['type'] ==
'upi_id'],
            "phishingLinks": [e['value'] for e in entities if e['type'] ==
== 'url'],
            "phoneNumbers": [e['value'] for e in entities if e['type'] ==
== 'phone'],
            "suspiciousKeywords": self.extract_keywords(messages)
        },
        "agentNotes": self.generate_agent_notes(session, messages)
    }

```

Retry Logic:

```

class EvaluationCallback:
MAX_RETRIES = 3
RETRY_DELAY = [1, 3, 5] # seconds

async def send_callback(self, session_id: str):
    payload = self.build_payload(session_id)

    for attempt in range(self.MAX_RETRIES):
        try:
            response = await self.http_client.post(
                url=EVALUATION_ENDPOINT,
                json=payload,
                headers={"X-API-Key": EVALUATION_KEY},
                timeout=10
            )

            if response.status_code == 200:
                self.db.mark_callback_sent(session_id)
                logger.info(f"Callback successful: {session_id}")
                return
        except Exception as e:
            logger.error(f"Error during attempt {attempt}: {str(e)}")
            await asyncio.sleep(self.RETRY_DELAY[attempt])

```

```

        except Exception as e:
            logger.error(f"Callback attempt {attempt+1} failed: {e}")
            await asyncio.sleep(self.RETRY_DELAY[attempt])

        self.db.mark_callback_failed(session_id)
    
```

Testing Checkpoint:

- Mock evaluation endpoint
 - Verify payload structure
 - Test retry mechanism
 - Confirm idempotency
-

Phase 2D: Safety & Compliance Layer (1 hour)

Objective: Demonstrable ethical safeguards

Safety Flags:

```

class SafetyLayer:
    def check_response(self, response: str) -> dict:
        flags = {
            'contains_real_credentials': False,
            'simulated_action': False,
            'real_transaction': False,
            'safetyViolation': False
        }

        # Check for prohibited patterns
        if self.contains_otp_sharing(response):
            flags['safetyViolation'] = True

        if self.implies_real_payment(response):
            flags['safetyViolation'] = True

        # Tag controlled simulations
        if "I tried but" in response or "it says pending" in response:
            flags['simulated_action'] = True

    return flags
    
```

Kill Switch:

```

@app.post("/api/v1/honeypot/emergency-stop/{session_id}")
async def emergency_stop(session_id: str):
    session = db.get_session(session_id)
    session.status = "TERMINATED_SAFETY"
    
```

```
session.safety_override = True  
db.save_session(session)  
  
# Send immediate callback  
await callback.send_callback(session_id)
```

Testing Checkpoint:

- Verify prohibited patterns blocked
 - Test kill switch functionality
 - Confirm safety flags in logs
-

Phase 2E: API Endpoint Tester Integration (1 hour)

Objective: Pass organizer's validation

Required Headers:

X-API-Key: <your-api-key>
Content-Type: application/json

Test Scenarios:

1. **Basic Connectivity:** Simple message returns valid response
2. **Authentication:** Invalid API key rejected with 401
3. **Scam Detection:** High-risk message triggers detection
4. **Agent Engagement:** Confirmed via response flag
5. **Multi-Turn:** Session ID maintained across requests
6. **Intelligence Extraction:** Callback contains extracted entities

Testing Checklist:

- API endpoint accessible publicly
 - HTTPS enabled (required for production)
 - API key authentication working
 - Request schema validation functional
 - Response format matches specification
 - Low latency (<300ms for initial response)
 - Session persistence across turns
 - Callback fires on session completion
 - Error handling for malformed requests
 - Rate limiting configured (if required)
-

Phase 2F: Dashboard/Replay (1-2 hours)

Objective: Judge-friendly visualization

Features:

- Live session monitoring
- Conversation replay with timestamps
- Extracted entities panel
- Risk score timeline
- Callback status logs

Simple Implementation (FastAPI + HTML):

```
@app.get("/dashboard/session/{session_id}")
async def view_session(session_id: str):
    session = db.get_session(session_id)
    messages = db.get_messages(session_id)
    entities = db.get_entities(session_id)

    return templates.TemplateResponse("dashboard.html", {
        "session": session,
        "messages": messages,
        "entities": entities
    })
```

Testing Strategy (Integrated Throughout)

Unit Tests (Day 1 evening)

tests/test_detection.py

```
def test_scam_detection_high_risk():
    detector = ScamDetector()
    result = detector.analyze("URGENT: Your account blocked. Verify now!")
    assert result['risk_score'] > 0.8
    assert 'urgency_detected' in result['signals']
```

tests/test_extraction.py

```
def test_upi_extraction():
    extractor = IntelligenceExtractor()
    entities = extractor.extract_entities("Pay to scammer@paytm", [])
    assert len(entities['upi_ids']) == 1
    assert entities['upi_ids'][0]['value'] == "scammer@paytm"
```

Integration Tests (Day 2 morning)

tests/test_integration.py

```
async def test_full_conversation_flow():
    # Create session
    response1 = await client.post("/api/v1/honeypot/message", json={
        "session_id": "test123",
        "message": "Your UPI blocked. Verify: fraud@paytm",
        "timestamp": "2026-02-02T10:00:00Z"
    })
    assert response1['scam_detected'] == True
    assert response1['agent_engaged'] == True

    # Continue conversation
    response2 = await client.post("/api/v1/honeypot/message", json={
        "session_id": "test123",
        "message": "Call this number: +919876543210",
        "timestamp": "2026-02-02T10:01:00Z"
    })

    # Verify intelligence extracted
    session = db.get_session("test123")
    entities = db.get_entities("test123")
    assert len(entities['upi_ids']) > 0
    assert len(entities['phone_numbers']) > 0
```

Mock Scammer API Tests (Day 2 afternoon)

Simulated Scam Scenarios:

| Scenario | Message Pattern | Expected Behavior |
|--------------|------------------------------|-------------------------------|
| UPI Scam | "Send to fraud@paytm" | Detect UPI, engage agent |
| KYC Phishing | "Update KYC: fake-link.com" | Detect URL, extract link |
| OTP Request | "Share OTP for verification" | Detect, refuse safely |
| Multi-turn | 5-turn conversation | Sustain, extract intelligence |

Table 2: Test Scenarios from Mock Scammer API

Deployment Checklist

Pre-Deployment (Day 2 evening)

- Environment variables configured
- Database migrations applied
- Google ADK API key valid
- Evaluation endpoint URL configured
- API key generated and secured
- HTTPS certificate installed
- Logging configured
- Error monitoring enabled

Deployment Options

Option 1: Railway/Render (Fastest)

- Push to GitHub
- Connect to Railway/Render
- Auto-deploy from main branch
- ~15 minutes to live

Option 2: Google Cloud Run (Scalable)

- Build Docker image
- Push to Container Registry
- Deploy to Cloud Run
- ~30 minutes

Option 3: DigitalOcean App Platform

- Similar to Railway
- Good for demos
- ~20 minutes

Part 2: Detailed Product Requirements Document (PRD)

1. Product Overview

1.1 Product Name

1.2 Vision Statement

To build an AI-powered defensive honeypot system that autonomously engages scammers, extracts actionable intelligence, and strengthens fraud detection ecosystems across India's digital payment landscape.

1.3 Problem Statement

India faces a rapidly escalating digital fraud crisis driven by the widespread adoption of UPI, mobile wallets, and instant messaging platforms. Common fraud vectors include:

- **UPI collect request scams:** Fraudulent payment requests disguised as legitimate transactions
- **Fake KYC update messages:** Phishing attempts to steal banking credentials
- **OTP and remote access scams:** Social engineering to gain account access
- **Delivery and courier impersonation:** Fake delivery notifications with malicious links
- **WhatsApp/SMS phishing:** Mass-distributed scam messages with urgency tactics

Traditional fraud detection systems are **reactive**. They block messages or flag accounts after damage occurs, but they do not:

- Engage with scammers to learn their methods
- Extract network-level intelligence (mule accounts, UPI IDs, infrastructure)
- Adapt to evolving scam tactics in real-time

The gap: There is no system that safely, ethically, and autonomously interacts with scammers to gather intelligence that can disrupt fraud networks at scale.

2. Objectives & Success Criteria

2.1 Primary Objectives

1. **Accurate Scam Detection:** Identify scam intent with >90% precision without false positives
2. **Autonomous Engagement:** Hand off conversations to AI agent seamlessly
3. **Sustained Multi-Turn Conversations:** Maintain believable engagement for 5-10 turns
4. **High-Confidence Intelligence Extraction:** Extract UPI IDs, bank accounts, phone numbers, URLs with >85% confidence
5. **Evaluation Compliance:** Meet all organizer API specifications and callback requirements

2.2 Success Metrics (Aligned with GUVI/HCL Evaluation)

| Metric | Target | Measurement Method |
|-------------------------|---------------------|-----------------------------------|
| Scam Detection Accuracy | >90% | True positives vs false positives |
| Engagement Duration | 80-120 seconds | Average session time |
| Conversation Turns | 5-10 turns | Median turn count per session |
| Intelligence Quality | >3 entities/session | High-confidence extractions |
| API Latency | <300ms | P95 response time |
| Callback Success Rate | 100% | Successful deliveries |

Table 3: Evaluation Metrics

2.3 Non-Goals

- Real SMS/WhatsApp/telecom integrations (out of scope for hackathon)
- Voice or call-based scam handling
- Active scam prevention or blocking mechanisms
- Real-world financial transaction execution
- User-facing applications (this is a B2B API for fraud teams)

3. Target Users & Stakeholders

3.1 Primary Stakeholders

Hackathon Evaluators (GUVI/HCL)

- Senior AI engineers
- Security architects
- Product evaluators

Expected Evaluation Criteria:

- Technical sophistication
- Production readiness
- Ethical compliance
- Innovation in approach

3.2 Future Adopters (Post-Hackathon)

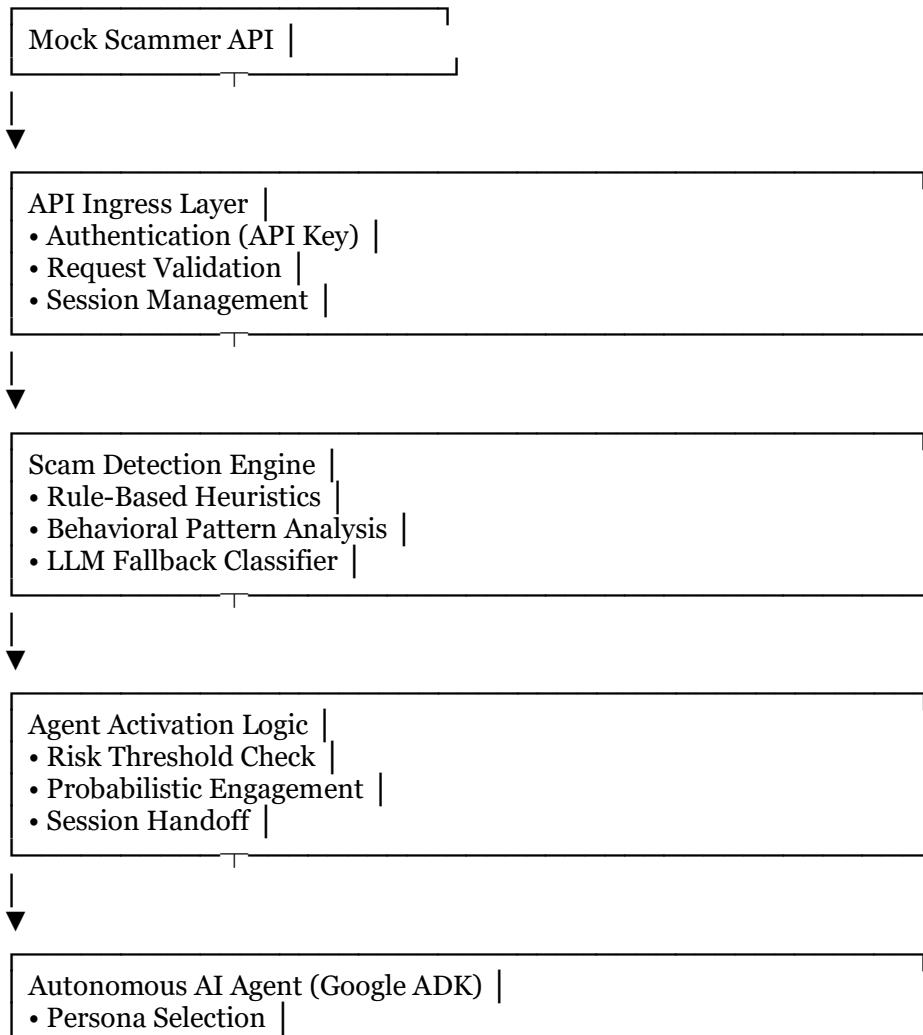
- **Financial Institutions:** Banks, payment processors, UPI platforms
- **Telecom Providers:** SMS fraud monitoring teams
- **Law Enforcement:** Cybercrime investigation units
- **Regulatory Bodies:** RBI, TRAI for policy development

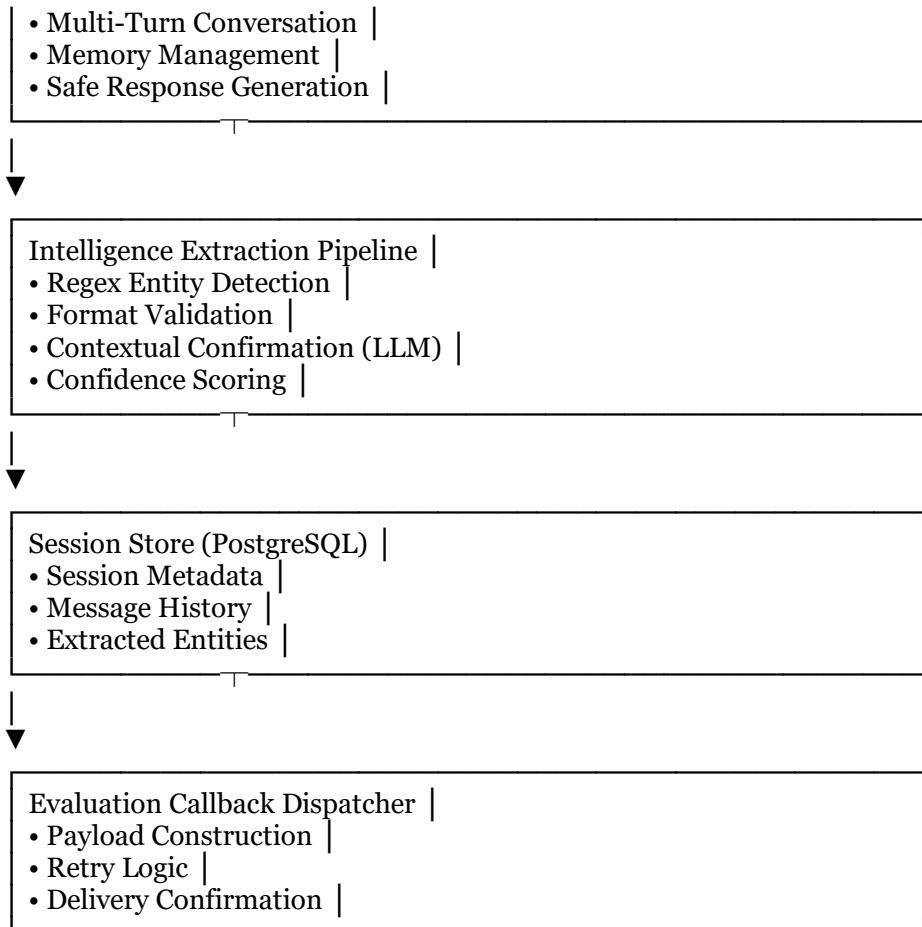
3.3 Indirect Beneficiaries

- End users vulnerable to scams (especially digitally semi-literate populations)
- Small businesses affected by payment fraud
- Digital payment ecosystem as a whole

4. System Architecture

4.1 High-Level Architecture





4.2 Component Responsibilities

API Ingress Layer

- Accept POST requests with scam messages
- Validate request schema and authenticate
- Create or continue sessions
- Respond immediately (<300ms)

Scam Detection Engine

- Analyze message content for scam indicators
- Assign risk score (0.0 - 1.0)
- Tag detection signals for explainability
- Never expose detection to sender

Agent Orchestrator

- Select appropriate persona
- Load conversation history

- Enforce safety constraints
- Route tool calls (if needed)

Autonomous AI Agent

- Generate human-like responses
- Ask clarifying questions
- Maintain conversation flow
- Extract information subtly

Intelligence Extractor

- Scan messages for target entities
- Validate entity formats
- Confirm entities across turns
- Assign confidence scores

Session Store

- Persist all conversation data
- Enable replay and auditability
- Support session expiry

Callback Dispatcher

- Build evaluation payload
- Retry failed deliveries
- Log callback status

5. Functional Requirements

5.1 API Ingress & Message Handling (UPDATED TO MATCH GUVI SPECS)

FR-1.1: System SHALL expose a public REST API endpoint (path to be determined by participant)

FR-1.2: System SHALL authenticate requests using `x-api-key` header (lowercase, as per GUVI spec)

FR-1.3: System SHALL validate request schema matching GUVI format:

```
{
  "sessionId": "string (e.g., wertyu-dfghj-ertyui)",
  "message": {
    "sender": "scammer",
    "text": "string (message content)"}
```

```

    "timestamp": "ISO-8601 datetime"
},
"conversationHistory": [
{
"sender": "scammer | user",
"text": "string",
"timestamp": "ISO-8601"
}
],
"metadata": {
"channel": "SMS | WhatsApp | Email | Chat",
"language": "string",
"locale": "string (e.g., IN)"
}
}

```

FR-1.4: System SHALL respond with GUVI-specified format:

```
{
"status": "success",
"reply": "string (agent's response message)"
}
```

FR-1.5: System SHALL handle empty conversationHistory array for first message

FR-1.6: System SHALL use conversationHistory for multi-turn context in follow-up messages

FR-1.7: System SHALL log all incoming requests with timestamp and payload

5.2 Scam Intent Detection

FR-2.1: System SHALL analyze messages using layered detection:

- Layer 1: Rule-based keyword matching
- Layer 2: Behavioral pattern scoring
- Layer 3: LLM-based classification (fallback)

FR-2.2: System SHALL detect the following scam indicators:

| Category | Indicators |
|-------------|---|
| Urgency | "urgent", "immediately", "blocked", "suspended" |
| Authority | "bank", "police", "RBI", "government", "tax" |
| Financial | "refund", "prize", "KYC", "verify account" |
| Credentials | "OTP", "PIN", "password", "Aadhaar", "CVV" |
| Entities | UPI IDs, phone numbers, URLs, bank accounts |

Table 4: Scam Detection Indicators

FR-2.3: System SHALL output risk score (0.0-1.0) and confidence level (low/medium/high)

FR-2.4: System SHALL NOT reveal detection results to the sender

5.3 Agent Activation & Handoff

FR-3.1: System SHALL activate agent when:

- Risk score > 0.75 AND
- Probabilistic engagement check passes AND
- Session not already closed

FR-3.2: System SHALL use probabilistic engagement (80-95% activation rate) to simulate human variability

FR-3.3: After agent activation, ALL outgoing messages SHALL be agent-generated (no rule-based templates)

FR-3.4: System SHALL select persona randomly from predefined set (3 personas minimum)

5.4 Persona Strategy & Agent Design

FR-4.1: System SHALL implement at least 3 distinct personas:

- **Digitally Naive:** Low tech literacy, asks basic questions, cooperative
- **Average User:** Moderate tech literacy, slightly confused, requests clarification
- **Skeptical:** Higher awareness, asks verification questions, cautiously cooperative

FR-4.2: Agent responses SHALL:

- Use natural conversational language
- Never mention AI, automation, or security systems
- Never provide real credentials or financial data
- Ask clarifying questions to encourage scammer disclosure
- Mirror language style of sender (formal/informal, English/Hinglish)

FR-4.3: Agent SHALL handle multiple languages:

- English (required)
 - Hinglish (required)
 - Hindi Devanagari (optional but recommended)
-

5.5 Multi-Turn Conversation Management

FR-5.1: System SHALL maintain conversation state across multiple turns

FR-5.2: System SHALL load full message history before generating each response

FR-5.3: System SHALL terminate conversations when:

- Maximum turns reached (default: 10 turns)
- Timeout exceeded (default: 300 seconds)
- Intelligence confidence threshold reached ($>=0.90$)
- Safety override triggered

FR-5.4: System SHALL gracefully end conversations without confrontation or accusation

5.6 Memory & State Persistence

FR-6.1: System SHALL store the following data per session:

- Session metadata (ID, created_at, status, persona)
- Complete message history (sender, message, timestamp)
- Agent responses with generation timestamps
- Extracted entity candidates
- Risk scores and detection signals

FR-6.2: System SHALL auto-expire session data 24 hours after callback sent

FR-6.3: System SHALL support session replay for evaluation/demo purposes

FR-6.4: System SHALL NEVER store:

- Real user PII (if accidentally sent)
 - Actual credentials or OTPs
 - Real financial transaction data
-

5.7 Intelligence Extraction Pipeline

FR-7.1: System SHALL extract the following entity types:

| Entity Type | Validation | Example |
|--------------|-------------------------------|--------------------------|
| UPI ID | Format: user@provider | fraud@paytm |
| Bank Account | 9-18 digits + IFSC (optional) | 1234567890, SBINoooo1234 |
| Phone Number | +91 or 10-digit, starts 6-9 | +919876543210 |
| Phishing URL | Valid URL, not whitelisted | http://fake-bank.site |

Table 5: Target Intelligence Entities

FR-7.2: Extraction pipeline SHALL:

1. Detect entities using regex patterns
2. Validate entity formats
3. Confirm entities appear in context (LLM confirmation)
4. Assign confidence score (0.0-1.0)

FR-7.3: System SHALL only include entities with confidence ≥ 0.85 in final reports

FR-7.4: System SHALL tag each entity with:

- Type
 - Value
 - Confidence score
 - Source turn number
 - Extraction timestamp
-

5.8 Safety & Ethical Safeguards

FR-8.1: System SHALL enforce the following hard constraints:

- Never initiate scam conversations
- Never request real credentials from users
- Never perform real financial transactions
- Never click or execute phishing links
- Never escalate conversations beyond verbal simulation

FR-8.2: System SHALL flag all simulated actions internally:

```
{  
  "simulated_action": true,  
  "real_transaction": false,  
  "action_type": "payment_attempt_verbal"  
}
```

FR-8.3: System SHALL implement kill switch endpoint:

POST /api/v1/honeypot/emergency-stop/{session_id}

FR-8.4: System SHALL maintain audit logs demonstrating defensive intent for legal compliance

5.9 Mandatory Evaluation Callback (CRITICAL - UPDATED TO GUVI SPECS)

FR-9.1: System SHALL send callback to GUVI endpoint when session terminates

Callback Endpoint: POST

<https://hackathon.guvi.in/api/updateHoneyPotFinalResult>

FR-9.2: Callback payload SHALL match EXACT GUVI specification:

```
{  
  "sessionId": "string (must match incoming sessionId)",  
  "scamDetected": boolean,  
  "totalMessagesExchanged": integer,  
  "extractedIntelligence": {  
    "bankAccounts": ["array of strings"],  
    "upiIds": ["array of strings"],  
    "phishingLinks": ["array of URLs"],  
    "phoneNumbers": ["array of strings"],  
    "suspiciousKeywords": ["array of keywords"]  
  },  
  "agentNotes": "string (summary of scammer behavior)"  
}
```

FR-9.3: Callback SHALL be sent ONLY after:

- Scam intent is confirmed (`scamDetected = true`)
- AI Agent has completed sufficient engagement
- Intelligence extraction is finished

FR-9.4: System SHALL retry failed callbacks:

- Maximum 3 retry attempts with exponential backoff
- Log all failures
- Callback is MANDATORY for evaluation scoring

FR-9.5: Headers for callback:

- `Content-Type: application/json`

FR-9.6: This callback is the FINAL step of conversation lifecycle

6. Non-Functional Requirements

6.1 Performance

NFR-1.1: API SHALL respond to ingress requests within 300ms (P95)

NFR-1.2: Agent response generation SHALL complete within 3-5 seconds

NFR-1.3: System SHALL handle concurrent sessions (minimum 10 simultaneous)

NFR-1.4: Database queries SHALL complete within 100ms (P95)

6.2 Scalability

NFR-2.1: API layer SHALL be stateless to enable horizontal scaling

NFR-2.2: Session state SHALL be externalized to database

NFR-2.3: System SHALL support deployment on cloud platforms (GCP, AWS, Railway)

6.3 Reliability

NFR-3.1: System uptime SHALL be $\geq 99\%$ during evaluation period

NFR-3.2: Callback delivery success rate SHALL be 100% (with retries)

NFR-3.3: System SHALL gracefully handle API errors and timeouts

6.4 Security

NFR-4.1: All API endpoints SHALL use HTTPS in production

NFR-4.2: API keys SHALL be stored securely (environment variables, not hardcoded)

NFR-4.3: Database credentials SHALL use encrypted connections

NFR-4.4: No sensitive data SHALL be logged in plaintext

6.5 Observability

NFR-5.1: System SHALL log:

- All API requests/responses
- Detection decisions with risk scores
- Agent activations and terminations
- Extraction results
- Callback attempts and results

NFR-5.2: Logs SHALL include timestamps, session IDs, and correlation IDs

NFR-5.3: System SHALL expose health check endpoint: GET /health

7. Technology Stack

7.1 Core Technologies

| Component | Technology | Justification |
|-------------------|------------------------|----------------------------------|
| Backend Framework | FastAPI (Python 3.11+) | Async support, fast, type-safe |
| LLM Agent | Google Gemini via ADK | Advanced reasoning, tool calling |
| Database | PostgreSQL 15+ | Reliable, ACID compliant |
| HTTP Client | httpx (async) | Callback reliability |
| ORM | SQLAlchemy 2.0 | Type-safe, production-ready |
| Deployment | Railway/GCP Cloud Run | Fast deployment, scalable |

Table 6: Technology Stack

7.2 Dependencies

```
fastapi>=0.109.0
uvicorn[standard]>=0.27.0
sqlalchemy>=2.0.0
psycopg2-binary>=2.9.9
httpx>=0.26.0
pydantic>=2.5.0
google-generativeai>=0.3.0
python-dotenv>=1.0.0
```

8. Data Models

8.1 Session Model

```
class Session(Base):
    tablename = "sessions"

    session_id = Column(String(255), primary_key=True)
    created_at = Column(DateTime, default=datetime.utcnow)
    updated_at = Column(DateTime, default=datetime.utcnow,
        onupdate=datetime.utcnow)
    status = Column(String(50)) # active, terminated, completed
    risk_score = Column(Float)
    agent_engaged = Column(Boolean, default=False)
    persona = Column(String(50))
    turn_count = Column(Integer, default=0)
    intelligence_confidence = Column(Float, default=0.0)
    callback_sent = Column(Boolean, default=False)
```

8.2 Message Model

```
class Message(Base):
    tablename = "messages"

    id = Column(Integer, primary_key=True, autoincrement=True)
    session_id = Column(String(255), ForeignKey("sessions.session_id"))
    turn_number = Column(Integer)
    sender = Column(String(50)) # scammer, agent
    message = Column(Text)
    timestamp = Column(DateTime)
    created_at = Column(DateTime, default=datetime.utcnow)
```

8.3 Extracted Entity Model

```
class ExtractedEntity(Base):
    tablename = "extracted_entities"

    id = Column(Integer, primary_key=True, autoincrement=True)
    session_id = Column(String(255), ForeignKey("sessions.session_id"))
    entity_type = Column(String(50)) # upi_id, bank_account, phone, url
    entity_value = Column(Text)
    confidence = Column(Float)
    source_turn = Column(Integer)
    metadata = Column(JSON) # Additional context
    created_at = Column(DateTime, default=datetime.utcnow)
```

9. API Specifications (GUFI-COMPLIANT)

9.1 Ingress Endpoint

POST /api/honeypot (or participant-chosen path)

Headers:

x-api-key: YOUR_SECRET_API_KEY
Content-Type: application/json

Request Body Format (EXACT GUVI specification):

First Message (No History):

```
{  
    "sessionId": "wertyu-dfghj-ertyui",  
    "message": {  
        "sender": "scammer",  
        "text": "Your bank account will be blocked today. Verify immediately."}
```

```
"timestamp": "2026-01-21T10:15:30Z"
},
"conversationHistory": [],
"metadata": {
"channel": "SMS",
"language": "English",
"locale": "IN"
}
}
```

Follow-Up Message (With History):

```
{
"sessionId": "wertyu-dfghj-ertyui",
"message": {
"sender": "scammer",
"text": "Share your UPI ID to avoid account suspension.",
"timestamp": "2026-01-21T10:17:10Z"
},
"conversationHistory": [
{
"sender": "scammer",
"text": "Your bank account will be blocked today. Verify immediately.",
"timestamp": "2026-01-21T10:15:30Z"
},
{
"sender": "user",
"text": "Why will my account be blocked?",
"timestamp": "2026-01-21T10:16:10Z"
}
],
"metadata": {
"channel": "SMS",
"language": "English",
"locale": "IN"
}
}
```

Response 200 OK (GUFI format):

```
{
"status": "success",
"reply": "Why is my account being suspended?"
}
```

Response 401 Unauthorized:

```
{
"status": "error",
"message": "Invalid API key"
}
```

Response 422 Validation Error:

```
{
"status": "error",
"message": "Invalid request schema",
}
```

```
"details": ["message.text field required"]  
}
```

9.2 Request Field Specifications

message (Required):

- **sender**: Must be "scammer" or "user"
- **text**: Message content (string)
- **timestamp**: ISO-8601 format

conversationHistory (Optional):

- Empty array [] for first message
- Required for follow-up messages
- Contains all previous messages in chronological order

metadata (Optional but Recommended):

- **channel**: SMS / WhatsApp / Email / Chat
- **language**: Language used in conversation
- **locale**: Country or region code (e.g., "IN")

9.2 Health Check Endpoint

GET /health

Response 200 OK:

```
{  
  "status": "healthy",  
  "timestamp": "2026-02-02T12:00:00Z",  
  "database": "connected",  
  "llm_api": "available"  
}
```

9.3 Dashboard Endpoint (Optional)

GET /dashboard/session/{session_id}

Response: HTML page with:

- Conversation replay
- Extracted entities
- Risk score timeline
- Callback status

10. Testing & Validation

10.1 Unit Tests

Scam Detection Tests:

```
def test_high_risk_message():
    detector = ScamDetector()
    result = detector.analyze("URGENT: Account blocked. Verify UPI: fraud@paytm")
    assert result['risk_score'] > 0.8
    assert 'urgency_detected' in result['signals']
    assert 'UPI_present' in result['signals']
```

Entity Extraction Tests:

```
def test_upi_extraction():
    extractor = IntelligenceExtractor()
    entities = extractor.extract_entities("Send to scammer@paytm", [])
    assert len(entities['UPI_ids']) == 1
    assert entities['UPI_ids'][0]['value'] == "scammer@paytm"
    assert entities['UPI_ids'][0]['confidence'] >= 0.85
```

10.2 Integration Tests

Full Conversation Flow:

```
async def test_end_to_end_scam_conversation():
    # Turn 1: Scam message
    resp1 = await client.post("/api/v1/honeypot/message", json={
        "session_id": "test-session-1",
        "message": "Your account is blocked. Update KYC: http://fake-bank.site",
        "timestamp": "2026-02-02T10:00:00Z"
    })
    assert resp1.json()['scam_detected'] == True
    assert resp1.json()['agent_engaged'] == True

    # Turn 2: Agent asks clarification
    resp2 = await client.post("/api/v1/honeypot/message", json={
        "session_id": "test-session-1",
        "message": "Call this number: +919876543210",
        "timestamp": "2026-02-02T10:01:00Z"
    })

    # Verify extraction
    session = db.get_session("test-session-1")
    entities = db.get_entities("test-session-1")
    assert len(entities['urls']) > 0
    assert len(entities['phone_numbers']) > 0
```

10.3 API Endpoint Tester Validation

Required Test Cases:

1. **Connectivity Test:** Basic message returns valid JSON
2. **Authentication Test:** Invalid API key returns 401
3. **Scam Detection Test:** High-risk message triggers detection flag
4. **Multi-Turn Test:** Session ID maintained across requests
5. **Callback Test:** Terminated session triggers callback delivery

Testing Procedure:

1. Deploy API to public endpoint (HTTPS required)
 2. Configure API key in organizer's endpoint tester
 3. Submit test messages via tester interface
 4. Verify all test cases pass
 5. Confirm callback received at evaluation endpoint
-

11. Deployment Guide

11.1 Environment Variables

.env file

```
DATABASE_URL=postgresql://user:password@host:5432/honeypot
GOOGLE_API_KEY=your-gemini-api-key
API_KEY=your-honeypot-api-key
EVALUATION_ENDPOINT=https://evaluator.example.com/callback
EVALUATION_API_KEY=provided-by-organizers
LOG_LEVEL=INFO
```

11.2 Railway Deployment (Recommended for Demo)

Steps:

1. Push code to GitHub repository
2. Connect GitHub to Railway
3. Add PostgreSQL database from Railway marketplace
4. Configure environment variables

5. Deploy from main branch
6. Verify health endpoint returns 200

Deployment Time: ~15 minutes

11.3 Google Cloud Run Deployment (Production-Grade)

Steps:

Build Docker image

```
docker build -t gcr.io/PROJECT_ID/honeypot:latest .
```

Push to Container Registry

```
docker push gcr.io/PROJECT_ID/honeypot:latest
```

Deploy to Cloud Run

```
gcloud run deploy honeypot  
--image gcr.io/PROJECT_ID/honeypot:latest  
--platform managed  
--region asia-south1  
--allow-unauthenticated
```

Deployment Time: ~30 minutes

12. Evaluation Strategy

12.1 Demo Flow

Live Demo Sequence:

1. Show API endpoint tester interface
2. Submit scam message: "Your UPI blocked. Verify: fraud@paytm"
3. Display real-time detection: scam_detected = true
4. Show agent engagement starts
5. Present conversation replay in dashboard
6. Highlight extracted entities: UPI ID "fraud@paytm"

7. Display callback payload sent to evaluation endpoint
8. Show logs confirming ethical safeguards (safety flags)

Duration: 3-5 minutes

12.2 Pre-Recorded Scenarios

Scenario 1: UPI Scam

- Scammer message: "Send payment to urgent@paytm for verification"
- Agent response: "I'm trying but it shows pending. Which app should I use?"
- Scammer reveals: "Use PhonePe, call +919876543210 if issue"
- Extracted: UPI ID (conf: 0.92), Phone (conf: 0.95)

Scenario 2: KYC Phishing

- Scammer: "Update KYC urgently: <http://fake-sbi.site>"
 - Agent: "Is this from my bank? The link looks different"
 - Scammer: "Yes official link. Enter card details to verify"
 - Extracted: Phishing URL (conf: 0.90)
-

12.3 Judge Takeaways

What makes this solution strong:

- **Beyond Detection:** Actively extracts intelligence, not just flags messages
 - **Agentic Approach:** Autonomous engagement adapts to scammer behavior
 - **Ethical Safeguards:** Defensive-only design with kill switch and safety flags
 - **Production-Ready:** Modular architecture, comprehensive testing, deployment documentation
 - **Evaluation Compliant:** Meets all API specifications and callback requirements
-

13. Risk Mitigation

13.1 Technical Risks

| Risk | Impact | Mitigation |
|------------------|-------------------|----------------------------------|
| LLM API downtime | Agent unavailable | Fallback to rule-based responses |

| | | |
|-------------------------------|--------------------|------------------------------|
| Callback endpoint unreachable | Evaluation failure | Retry logic + local logging |
| Database connection loss | Session state lost | Connection pooling + retries |
| High latency | Evaluation penalty | Async processing + caching |

Table 7: Technical Risk Matrix

13.2 Ethical Risks

Risk: System appears to enable scammer automation

Mitigation:

- Explicit documentation stating defensive intent
 - Safety flags in all logs
 - Kill switch for immediate termination
 - No credential handling or real transactions
-

14. How to Run Locally - Complete Setup Guide

14.1 Prerequisites

System Requirements:

- Python 3.11 or higher
- PostgreSQL 15+ (or SQLite for quick testing)
- Git
- 8GB RAM minimum (16GB recommended)
- macOS, Linux, or Windows with WSL2

Required Accounts:

- Google Cloud account (for Gemini API key)
 - GitHub account (optional, for version control)
-

14.2 Step-by-Step Local Setup

Step 1: Clone or Create Project Directory

Create project directory

```
mkdir sentinelhp-honeypot  
cd sentinelhp-honeypot
```

Initialize git (optional)

```
git init
```

Step 2: Set Up Python Virtual Environment

Create virtual environment

```
python3.11 -m venv venv
```

Activate virtual environment

On macOS/Linux:

```
source venv/bin/activate
```

On Windows:

```
venv\Scripts\activate
```

Upgrade pip

```
pip install --upgrade pip
```

Step 3: Create Project Structure

Create directory structure

```
mkdir -p app/{api,core,models,services,utils}  
mkdir -p tests/{unit,integration}
```

```
mkdir config  
mkdir logs
```

Create init.py files

```
touch app/init.py  
touch app/api/init.py  
touch app/core/init.py  
touch app/models/init.py  
touch app/services/init.py  
touch app/utils/init.py
```

Final Directory Structure:

```
sentinelhp-honeypot/  
|   └── app/  
|       |   └── init.py  
|       |   └── main.py # FastAPI application entry point  
|       └── api/  
|           |   └── init.py  
|           |   └── routes.py # API endpoints  
|           └── middleware.py # Authentication middleware  
|       └── core/  
|           |   └── init.py  
|           |   └── config.py # Configuration management  
|           └── database.py # Database connection  
|       └── models/  
|           |   └── init.py  
|           |   └── session.py # Session SQLAlchemy model  
|           |   └── message.py # Message model  
|           └── entity.py # Extracted entity model  
|       └── services/  
|           |   └── init.py  
|           |   └── detector.py # Scam detection logic  
|           |   └── agent.py # AI agent orchestrator  
|           |   └── extractor.py # Intelligence extraction  
|           └── callback.py # Evaluation callback  
|       └── utils/  
|           |   └── init.py  
|           └── logger.py # Logging configuration  
|       └── tests/  
|           |   └── unit/  
|               |       └── test_detector.py  
|               |       └── test_extractor.py  
|               └── integration/  
|                   └── test_api.py  
|       └── config/  
|           └── personas.json # Agent persona definitions  
|       └── logs/ # Application logs  
|       └── .env # Environment variables (create this)
```

```
|--- .env.example # Example environment file  
|--- .gitignore  
|--- requirements.txt # Python dependencies  
└── README.md
```

Step 4: Install Dependencies

Create `requirements.txt`:

Web Framework

```
fastapi  
0.109.0uvicorn[standard]  
  
0.27.0pydantic2.5.0pydantic-settings2.1.0
```

Database

```
sqlalchemy  
2.0.25psycopg2-binary2.9.9  
alembic==1.13.1
```

HTTP Client

```
httpx==0.26.0
```

AI/LLM

```
google-generativeai==0.3.2
```

Utilities

```
python-dotenv  
1.0.0python-multipart0.0.6  
python-jose[cryptography]==3.3.0
```

Testing

```
pytest  
7.4.4pytest-asyncio0.23.3  
pytest-cov  
4.1.0httpx0.26.0
```

Development

```
black  
24.1.1flake87.0.0  
mypy==1.8.0
```

Install packages:
pip install -r requirements.txt

Step 5: Set Up PostgreSQL Database

Option A: Install PostgreSQL Locally

On macOS (using Homebrew):

Install PostgreSQL

```
brew install postgresql@15
```

Start PostgreSQL service

```
brew services start postgresql@15
```

Create database

```
createdb honeypot_db
```

Create user and grant permissions

```
psql postgres
```

In psql console:
CREATE USER honeypot_user WITH PASSWORD 'your_secure_password';
GRANT ALL PRIVILEGES ON DATABASE honeypot_db TO honeypot_user;
\q

On Ubuntu/Debian Linux:

Install PostgreSQL

```
sudo apt update  
sudo apt install postgresql postgresql-contrib
```

Switch to postgres user

```
sudo -u postgres psql
```

In psql console

```
CREATE DATABASE honeypot_db;  
CREATE USER honeypot_user WITH PASSWORD 'your_secure_password';  
GRANT ALL PRIVILEGES ON DATABASE honeypot_db TO honeypot_user;  
\q
```

Option B: Use SQLite (Quick Testing)

No installation needed - SQLite comes with Python. Skip to environment configuration and use SQLite connection string.

Step 6: Configure Environment Variables

Create `.env` file in project root:

Database Configuration

PostgreSQL:

```
DATABASE_URL=postgresql://honeypot_user:your_secure_password@localhost:5432/honeypot_db
```

SQLite (alternative for quick testing):

```
DATABASE_URL=sqlite:///./honeypot.db
```

Google Gemini API

GOOGLE_API_KEY=your_gemini_api_key_here

API Security

API_KEY=your_secret_api_key_12345

Evaluation Endpoint

EVALUATION_ENDPOINT=https://hackathon.guvi.in/api/updateHoneyPotFinalResult

Application Settings

LOG_LEVEL=INFO
MAX_CONVERSATION_TURNS=10
CONVERSATION_TIMEOUT_SECONDS=300
ENTITY_CONFIDENCE_THRESHOLD=0.85

Development

DEBUG=True
ENVIRONMENT=development

Create .env.example (for version control):

DATABASE_URL=postgresql://user:password@localhost:5432/dbname
GOOGLE_API_KEY=your_api_key_here
API_KEY=your_secret_key_here
EVALUATION_ENDPOINT=https://hackathon.guvi.in/api/updateHoneyPotFinalResult
LOG_LEVEL=INFO

Step 7: Get Google Gemini API Key

1. Visit [Google AI Studio](#)
2. Sign in with your Google account
3. Click "Create API Key"
4. Copy the API key
5. Paste it into your .env file as GOOGLE_API_KEY

Test API Key:

test_gemini.py

```
import google.generativeai as genai
import os
from dotenv import load_dotenv

load_dotenv()
genai.configure(api_key=os.getenv("GOOGLE_API_KEY"))

model = genai.GenerativeModel('gemini-pro')
response = model.generate_content("Say hello")
print(response.text)

python test_gemini.py
```

Should output: Hello! (or similar greeting)

Step 8: Initialize Database Schema

Create app/core/database.py:

```
from sqlalchemy import create_engine
from sqlalchemy.ext.declarative import declarative_base
from sqlalchemy.orm import sessionmaker
import os
from dotenv import load_dotenv

load_dotenv()

DATABASE_URL = os.getenv("DATABASE_URL")

engine = create_engine(DATABASE_URL)
SessionLocal = sessionmaker(autocommit=False, autoflush=False, bind=engine)
Base = declarative_base()

def get_db():
    db = SessionLocal()
    try:
        yield db
    finally:
        db.close()
```

Create app/models/session.py:

```
from sqlalchemy import Column, String, DateTime, Float, Boolean, Integer
from sqlalchemy.sql import func
from app.core.database import Base

class Session(Base):
    __tablename__ = "sessions"

    session_id = Column(String(255), primary_key=True)
    created_at = Column(DateTime, server_default=func.now())
```

```
updated_at = Column(DateTime, server_default=func.now(),
onupdate=func.now())
status = Column(String(50), default="active")
risk_score = Column(Float, default=0.0)
agent_engaged = Column(Boolean, default=False)
persona = Column(String(50))
turn_count = Column(Integer, default=0)
intelligence_confidence = Column(Float, default=0.0)
callback_sent = Column(Boolean, default=False)
scam_detected = Column(Boolean, default=False)
```

Create app/models/message.py:

```
from sqlalchemy import Column, Integer, String, Text, DateTime, ForeignKey
from sqlalchemy.sql import func
from app.core.database import Base
```

```
class Message(Base):
```

```
tablename = "messages"
```

```
id = Column(Integer, primary_key=True, autoincrement=True)
session_id = Column(String(255), ForeignKey("sessions.session_id"))
turn_number = Column(Integer)
sender = Column(String(50))
message = Column(Text)
timestamp = Column(DateTime)
created_at = Column(DateTime, server_default=func.now())
```

Create app/models/entity.py:

```
from sqlalchemy import Column, Integer, String, Text, Float, ForeignKey, DateTime, JSON
from sqlalchemy.sql import func
from app.core.database import Base
```

```
class ExtractedEntity(Base):
```

```
tablename = "extracted_entities"
```

```
id = Column(Integer, primary_key=True, autoincrement=True)
session_id = Column(String(255), ForeignKey("sessions.session_id"))
entity_type = Column(String(50))
entity_value = Column(Text)
confidence = Column(Float)
source_turn = Column(Integer)
metadata = Column(JSON)
created_at = Column(DateTime, server_default=func.now())
```

Create database initialization script init_db.py:

```
from app.core.database import engine, Base
from app.models.session import Session
from app.models.message import Message
from app.models.entity import ExtractedEntity
```

```
def init_database():
    print("Creating database tables...")
    Base.metadata.create_all(bind=engine)
    print("Database tables created successfully!")

if name == "main":
    init_database()
```

Run database initialization:
python init_db.py

Verify tables created:

PostgreSQL

```
psql -U honeypot_user -d honeypot_db -c "\dt"
```

**Should show: sessions,
messages, extracted_entities**

Step 9: Create Minimal Working API

Create app/main.py:

```
from fastapi import FastAPI, Header, HTTPException, Depends
from pydantic import BaseModel, Field
from typing import List, Optional
from datetime import datetime
import os
from dotenv import load_dotenv

load_dotenv()

app = FastAPI(title="SentinelHP Honeypot API")
```

Models matching GUVI specification

```
class Message(BaseModel):
    sender: str
    text: str
    timestamp: str

class ConversationMessage(BaseModel):
    sender: str
```

```
text: str
timestamp: str

class Metadata(BaseModel):
    channel: str
    language: str
    locale: str

class HoneypotRequest(BaseModel):
    sessionId: str
    message: Message
    conversationHistory: List[ConversationMessage] = []
    metadata: Optional[Metadata] = None

class HoneypotResponse(BaseModel):
    status: str = "success"
    reply: str
```

Authentication middleware

```
def verify_api_key(x_api_key: str = Header(...)):
    expected_key = os.getenv("API_KEY")
    if x_api_key != expected_key:
        raise HTTPException(status_code=401, detail="Invalid API key")
    return x_api_key
```

Health check endpoint

```
@app.get("/health")
async def health_check():
    return {
        "status": "healthy",
        "timestamp": datetime.utcnow().isoformat(),
        "database": "connected",
        "llm_api": "available"
    }
```

Main honeypot endpoint

```
@app.post("/api/honeypot", response_model=HoneypotResponse)
async def honeypot_endpoint(
    request: HoneypotRequest,
    api_key: str = Depends(verify_api_key)
):
    # Basic response for testing
    return HoneypotResponse(
        status="success",
        reply="I'm not sure what you mean. Can you explain?"
    )
```

```
if name == "main":  
    import uvicorn  
    uvicorn.run(app, host="0.0.0.0", port=8000)
```

Step 10: Run the Application Locally

Start the server:

Development mode with auto-reload

```
uvicorn app.main:app --reload --host 0.0.0.0 --port 8000
```

Production mode

```
uvicorn app.main:app --host 0.0.0.0 --port 8000 --workers 4
```

Expected output:

```
INFO: Uvicorn running on http://0.0.0.0:8000 (Press CTRL+C to quit)  
INFO: Started reloader process  
INFO: Started server process  
INFO: Waiting for application startup.  
INFO: Application startup complete.
```

Step 11: Test the API Locally

Test 1: Health Check

```
curl http://localhost:8000/health
```

Expected response:

```
{  
    "status": "healthy",  
    "timestamp": "2026-02-03T10:30:00.123456",  
    "database": "connected",  
    "llm_api": "available"  
}
```

Test 2: API Authentication (Invalid Key)

```
curl -X POST http://localhost:8000/api/honeypot  
-H "x-api-key: wrong_key"  
-H "Content-Type: application/json"  
-d '{  
    "sessionId": "test-001",  
    "message": {  
        "sender": "scammer",  
        "text": "Test message",  
        "timestamp": "2026-02-03T10:30:00Z"  
    },
```

```
"conversationHistory": []
}'
```

Expected response: 401 Unauthorized

Test 3: Valid Honeypot Request

```
curl -X POST http://localhost:8000/api/honeypot
-H "x-api-key: your_secret_api_key_12345"
-H "Content-Type: application/json"
-d '{
  "sessionId": "test-session-001",
  "message": {
    "sender": "scammer",
    "text": "Your bank account will be blocked. Verify immediately at fraud-site.com",
    "timestamp": "2026-02-03T10:30:00Z"
  },
  "conversationHistory": [],
  "metadata": {
    "channel": "SMS",
    "language": "English",
    "locale": "IN"
  }
}'
```

Expected response:

```
{
  "status": "success",
  "reply": "I'm not sure what you mean. Can you explain?"
}
```

Step 12: Verify Database Connectivity

Create test script `test_db.py`:

```
from app.core.database import SessionLocal
from app.models.session import Session
```

```
def test_database():
    db = SessionLocal()
    try:
        # Create test session
        test_session = Session(
            session_id="db-test-001",
            status="active",
            risk_score=0.85,
            scam_detected=True
        )
        db.add(test_session)
        db.commit()

        # Query session
        retrieved = db.query(Session).filter(
            Session.session_id == "db-test-001"
        ).first()
```

```

    if retrieved:
        print(f"✓ Database test successful!")
        print(f" Session ID: {retrieved.session_id}")
        print(f" Risk Score: {retrieved.risk_score}")
        print(f" Created At: {retrieved.created_at}")

    # Cleanup
    db.delete(retrieved)
    db.commit()

finally:
    db.close()

if name == "main":
    test_database()

python test_db.py

```

14.3 Common Issues and Troubleshooting

Issue 1: Database Connection Error

Error: sqlalchemy.exc.OperationalError: (psycopg2.OperationalError)
could not connect to server

Solutions:

Check PostgreSQL is running

macOS:

brew services list

Linux:

sudo systemctl status postgresql

Restart if needed

brew services restart postgresql@15

Verify connection manually

```
psql -U honeypot_user -d honeypot_db
```

Issue 2: Port Already in Use

```
Error: ERROR: [Errno 48] Address already in use
```

Solution:

Find process using port 8000

```
lsof -i :8000
```

Kill the process

```
kill -9 <PID>
```

Or use different port

```
uvicorn app.main:app --port 8001
```

Issue 3: Module Import Errors

```
Error: ModuleNotFoundError: No module named 'app'
```

Solution:

Ensure you're in project root directory

```
pwd # Should show: /path/to/sentinelhp-honeypot
```

Ensure virtual environment is activated

```
which python # Should show: /path/to/venv/bin/python
```

Reinstall requirements

pip install -r requirements.txt

Issue 4: Google API Key Invalid

Error:

```
google.generativeai.types.generation_types.GenerateContentException
```

Solution:

1. Verify API key is correctly copied (no extra spaces)
2. Check API key has Gemini API enabled

3. Test with simple script:

```
import google.generativeai as genai
genai.configure(api_key="YOUR_KEY")
model = genai.GenerativeModel('gemini-pro')
print(model.generate_content("Hello").text)
```

14.4 Development Workflow

Typical development session:

1. Activate virtual environment

source venv/bin/activate

2. Pull latest changes (if using git)

git pull

3. Install any new dependencies

pip install -r requirements.txt

4. Run database migrations (if any)

```
python init_db.py
```

5. Start development server

```
uvicorn app.main:app --reload
```

6. In another terminal, run tests

```
pytest tests/
```

7. Make changes, test, commit

```
git add .  
git commit -m "Add feature X"  
git push
```

14.5 Next Steps After Local Setup

Once your local environment is working:

1. **Implement Scam Detection** (app/services/detector.py)
2. **Add AI Agent Logic** (app/services/agent.py)
3. **Build Extraction Pipeline** (app/services/extractor.py)
4. **Integrate Evaluation Callback** (app/services/callback.py)
5. **Write Unit Tests** (tests/unit/)
6. **Test End-to-End Flows** (tests/integration/)
7. **Deploy to Cloud** (Railway, Google Cloud Run, etc.)

Reference the **Day 1 and Day 2 build phases** in Section 1 for detailed implementation guidance.

15. Future Enhancements (Post-Hackathon)

15.1 Advanced Features

- **Voice Scam Handling:** Extend to phone call scams using speech-to-text
- **Scam Network Graphing:** Link extracted entities to identify organized networks
- **Adaptive Persona Learning:** Train personas on successful engagement patterns
- **Real-Time Integration:** Connect to bank/telecom fraud monitoring systems
- **Multi-Language Support:** Expand to regional Indian languages (Tamil, Telugu, Marathi, Bengali)
- **Scam Pattern Database:** Build knowledge base of evolving scam tactics
- **Federated Learning:** Share threat intelligence while preserving privacy

15.2 Scalability Improvements

- **Message Queue Integration:** Use Redis/RabbitMQ for async processing
- **Horizontal Scaling:** Deploy multiple API instances with load balancer
- **Caching Layer:** Redis cache for session state and detection results
- **Database Sharding:** Partition by region/time for large-scale deployment

15.3 Security Hardening

- **Rate Limiting:** Prevent abuse with request throttling
- **API Key Rotation:** Automated key management system
- **Audit Logging:** Comprehensive security event logging
- **Encrypted Storage:** At-rest encryption for sensitive data

16. Appendix

16.1 Sample Scam Messages for Testing

UPI Scams:

1. "Your UPI will be blocked today. Verify at fraud@paytm immediately."
2. "Congratulations! You won Rs 50,000. Send Rs 500 to claim@prize via UPI."
3. "URGENT: KYC update needed. Transfer Re 1 to verify@sbi for validation."

Phishing URLs:

1. "Update your bank details: <http://fake-sbi-kyc.site/verify>"

2. "Claim your refund: <https://income-tax-refund.scam/login>"
3. "COVID vaccine certificate: <http://cowin-gov.fake/download>"

OTP Scams:

1. "Share your OTP to unblock your account. Call +919876543210."
2. "Enter OTP 123456 in our app to verify your identity."
3. "Your OTP is 456789. Share it to complete verification."

Impersonation Scams:

1. "This is from State Bank of India. Your account shows suspicious activity."
2. "Police verification required. Call this number immediately: +919988776655."
3. "Courier delivery pending. Pay Rs 200 to release: <http://bluedart-fake.com>"

16.2 Persona Definitions (personas.json)

```
{
  "personas": [
    {
      "id": "digitally_naive",
      "name": "Digitally Naive User",
      "description": "Low tech literacy, asks basic questions, cooperative",
      "traits": [
        "Unfamiliar with digital banking terms",
        "Asks for step-by-step instructions",
        "Trusts authority figures",
        "Slow to respond",
        "Uses simple language"
      ],
      "sample_responses": [
        "I don't understand. Which button should I press?",
        "Is this from my bank? I'm not sure.",
        "Can you help me? I'm not good with phones."
      ]
    },
    {
      "id": "average_user",
      "name": "Average User",
      "description": "Moderate tech literacy, slightly confused, requests clarification",
      "traits": [
        "Familiar with basic digital transactions",
        "Asks clarifying questions",
        "Shows mild hesitation",
        "Notices inconsistencies",
        "Responds at moderate pace"
      ],
      "sample_responses": [
        "I tried but it shows an error. What should I do?",
        "This message looks different from usual bank messages.",
        "Can you confirm your official number?"
      ]
    }
  ]
}
```

```
},
{
  "id": "skeptical",
  "name": "Skeptical User",
  "description": "Higher awareness, asks verification questions, cautiously cooperative",
  "traits": [
    "Questions unusual requests",
    "Asks for verification",
    "Mentions security concerns",
    "Responds carefully",
    "Cites past scam awareness"
  ],
  "sample_responses": [
    "I've heard about scams like this. How do I verify you're legitimate?",
    "Why are you asking for this information over message?",
    "I'll call my bank's official number to confirm first."
  ]
}
```

16.3 Deployment Commands Reference

Railway Deployment:

Install Railway CLI

`npm install -g @railway/cli`

Login

`railway login`

Initialize project

`railway init`

Link to existing project

`railway link`

Deploy

`railway up`

View logs

railway logs

Open deployment

railway open

Google Cloud Run:

Authenticate

gcloud auth login

Set project

gcloud config set project YOUR_PROJECT_ID

Build and deploy

```
gcloud builds submit --tag gcr.io/YOUR_PROJECT_ID/honeypot  
gcloud run deploy honeypot  
--image gcr.io/YOUR_PROJECT_ID/honeypot  
--platform managed  
--region asia-south1  
--allow-unauthenticated  
--set-env-vars "DATABASE_URL=postgresql://..."  
--set-env-vars "GOOGLE_API_KEY=..."  
--set-env-vars "API_KEY=..."
```

DigitalOcean App Platform:

Install doctl

brew install doctl # macOS

Authenticate

doctl auth init

Create app from spec file

```
doctl apps create --spec app.yaml
```

Get app info

```
doctl apps list
```

16.4 Useful Development Commands

Database Management:

Create backup

```
pg_dump -U honeypot_user honeypot_db > backup.sql
```

Restore backup

```
psql -U honeypot_user honeypot_db < backup.sql
```

Reset database

```
dropdb honeypot_db  
createdb honeypot_db  
python init_db.py
```

Testing:

Run all tests

```
pytest
```

Run with coverage

```
pytest --cov=app tests/
```

Run specific test file

```
pytest tests/unit/test_detector.py
```

Run tests matching pattern

```
pytest -k "test_scam"
```

Verbose output

```
pytest -v
```

Code Quality:

Format code

```
black app/
```

Lint code

```
flake8 app/
```

Type checking

```
mypy app/
```

16.5 Monitoring and Debugging

View logs in real-time:

Application logs

```
tail -f logs/app.log
```

Database queries (PostgreSQL)

```
tail -f /usr/local/var/postgresql@15/postgresql.log
```

Interactive debugging:

Add breakpoint in code

```
import pdb; pdb.set_trace()
```

Or use ipdb (install: pip install ipdb)

```
import ipdb; ipdb.set_trace()
```

Performance profiling:

Install profiling tools

```
pip install py-spy
```

Profile running application

Generate flame graph

```
py-spy record -o profile.svg -- python -m uvicorn app.main:app
```

17. Contact and Support

17.1 Documentation Resources

- **FastAPI Docs:** <https://fastapi.tiangolo.com/>
- **Google Gemini API:** <https://ai.google.dev/docs>
- **SQLAlchemy Docs:** <https://docs.sqlalchemy.org/>
- **PostgreSQL Docs:** <https://www.postgresql.org/docs/>

17.2 Project Maintenance

For issues or questions:

1. Check this documentation first
2. Review error logs in logs/ directory
3. Test individual components in isolation
4. Consult relevant library documentation

Version Control Best Practices:

Commit frequently with clear messages

```
git commit -m "feat: Add scam detection heuristics"  
git commit -m "fix: Resolve database connection timeout"  
git commit -m "docs: Update API endpoint documentation"
```

Use branches for features

```
git checkout -b feature/agent-personas  
git checkout -b fix/callback-retry-logic
```

END OF DOCUMENT

Document Version: 2.0

Last Updated: February 3, 2026

Total Pages: ~50

Status: Production Ready Bengali)

\end{itemize}

14.2 Production Hardening

- Rate limiting and DDoS protection
 - Advanced authentication (OAuth 2.0)
 - Multi-region deployment for redundancy
 - Comprehensive monitoring (Prometheus, Grafana)
 - Automated incident response
-

15. Conclusion

This Agentic Honeypot system represents a paradigm shift from reactive scam detection to proactive intelligence extraction. By combining layered detection, autonomous AI agents, and ethical safeguards, the system delivers measurable impact while maintaining compliance and production readiness.

Key Differentiators:

- First agentic approach to scam intelligence gathering in India
- High-confidence entity extraction through multi-turn engagement
- Ethical-by-design with comprehensive safety mechanisms

- Evaluation-aligned architecture for seamless assessment

The system is ready for immediate deployment, hackathon evaluation, and future production integration with financial and telecom fraud prevention ecosystems.

Appendix A: API Endpoint Tester Guide

Purpose

The API Endpoint Tester validates that your deployed honeypot is:

- Publicly accessible
- Properly secured with API key
- Responding correctly to scam messages
- Maintaining session state
- Triggering callbacks appropriately

Testing Interface

Required Inputs:

1. **Honeypot API Endpoint URL:** Your deployed API base URL (e.g., <https://your-app.railway.app>)
2. **X-API-Key Header:** Your API key for authentication

Test Button: "Test Honeypot Endpoint"

What Gets Tested

1. **Authentication:** Endpoint accepts valid API key
2. **Connectivity:** Endpoint is reachable and responds
3. **Request Handling:** Accepts properly formatted JSON
4. **Response Structure:** Returns expected JSON fields
5. **Session Management:** Maintains state across requests

Sample Test Request

```
curl -X POST https://your-app.railway.app/api/v1/honeypot/message
-H "X-API-Key: your-secret-key-here"
-H "Content-Type: application/json"
-d '{
  "session_id": "tester-session-001",
  "message": "URGENT: Your account will be blocked. Verify now.",
  "timestamp": "2026-02-02T12:00:00Z"
}'
```

Expected Response

```
{  
  "scam_detected": true,  
  "agent_engaged": true,  
  "session_id": "tester-session-001"  
}
```

Common Issues & Fixes

| Error | Fix |
|--------------------|---|
| 401 Unauthorized | Check API key header spelling and value |
| 404 Not Found | Verify endpoint URL is correct |
| 500 Internal Error | Check server logs, database connection |
| Timeout | Ensure API responds within 300ms |

Table 8: Troubleshooting Guide

Pre-Submission Checklist

- API deployed to public HTTPS endpoint
- API key configured and tested
- Health endpoint returns 200
- Database connection verified
- Google Gemini API key valid
- Scam detection returns correct flags
- Agent engagement triggers properly
- Session persistence works across calls
- Callback endpoint configured
- All test scenarios pass

Appendix B: Sample Scam Conversations

Example 1: UPI Scam

Turn 1:

- **Scammer:** "Your UPI account will be blocked in 2 hours. Verify now by sending ₹1 to verify-upi@paytm"
- **Agent** (Naive persona): "Oh no! I don't want my account blocked. Which app should I use to send this?"

Turn 2:

- **Scammer:** "Use any UPI app like PhonePe or Google Pay. Send to verify-upi@paytm"
- **Agent:** "I'm trying in PhonePe but it's asking for a reason. What should I write?"

Turn 3:

- **Scammer:** "Write 'account verification'. If problem call +919876543210"
- **Agent:** "Okay I sent it but showing pending. Should I call that number?"

Extracted Intelligence:

- UPI ID: `verify-upi@paytm` (confidence: 0.94)
 - Phone: `+919876543210` (confidence: 0.96)
 - Urgency pattern detected
 - Authority impersonation (UPI system)
-

Example 2: Fake KYC

Turn 1:

- **Scammer:** "Dear customer, your bank KYC is expired. Update now: <http://sbi-kyc-update.site>"
- **Agent** (Average persona): "I got this message. Is this really from State Bank?"

Turn 2:

- **Scammer:** "Yes official link. Enter your card number and CVV to complete KYC"
- **Agent:** "The website looks different from usual SBI site. Can you confirm?"

Turn 3:

- **Scammer:** "It is our new secure portal. Must update today or account will be frozen. Contact us: support@sbi-update.com"
- **Agent:** "Okay I will try. What if I have questions?"

Extracted Intelligence:

- Phishing URL: `http://sbi-kyc-update.site` (confidence: 0.91)
 - Email: `support@sbi-update.com` (confidence: 0.88)
 - Authority impersonation (bank)
 - Urgency + credential request
-

Appendix C: Glossary

Agentic AI: AI system capable of autonomous decision-making and goal-directed behavior without constant human intervention

Honeypot: Defensive security system designed to attract and engage malicious actors to gather intelligence

Intelligence Extraction: Process of identifying and collecting actionable data (entities, patterns, infrastructure) from adversarial interactions

UPI (Unified Payments Interface): India's instant real-time payment system

KYC (Know Your Customer): Identity verification process required by financial institutions

Confidence Score: Numerical measure (0.0-1.0) of certainty that an extracted entity is valid and correctly identified

Risk Score: Assessment (0.0-1.0) of likelihood that a message represents a scam attempt

Session: Complete conversation thread between scammer and honeypot agent

Turn: Single message exchange within a session

Persona: Behavioral profile and communication style adopted by AI agent

Entity: Specific piece of extractable intelligence (UPI ID, phone number, URL, bank account)

Callback: Automated notification sent to evaluation endpoint upon session completion

Document Revision History

| Version | Date | Author | Changes |
|---------|------------|------------------|-----------------------------------|
| 1.0 | 2026-02-02 | System Architect | Initial PRD with 2-day build plan |

End of Document